

CS 480 – INTRODUCTION TO ARTIFICIAL INTELLIGENCE

TOPIC: LEARNING – INTRO

Mustafa Bilgic



<http://www.cs.iit.edu/~mbilgic>



<https://twitter.com/bilgicm>

LEARNING

- What's learning?
- Intro to Chapter 19: *“In which we describe agents that can improve their behavior through diligent study of past experiences and predictions about the future.”*
- We do not make any philosophical statements about whether the agent is *truly* learning
- *“An agent is learning if it improves its performance after making observations about the world.”*

WHY LEARN AND NOT PROGRAM DIRECTLY?

- We cannot anticipate all possible situations that the agent might find itself in
- Time/location/context changes knowledge and rules
- We might not know the solution crisp enough to program it
- We might not have time to encode all the knowledge

WHAT TO LEARN?

- Which action to take in a state (state \rightarrow action)
- Outcomes of our actions (action \rightarrow state)
- Mapping percepts to world states (percept \rightarrow state)
- Utility of the states (state \rightarrow utility)
- and more...

FEEDBACK

1. Unsupervised learning

- No feedback; the agent discovers patterns in the data
- E.g., clustering, dimensionality reduction, outlier detection

2. Supervised learning

- Feedback: input-output pairs
- E.g., classification, regression, ranking

3. Reinforcement learning

- Feedback: rewards

EPISODIC VS SEQUENTIAL

- Supervised and unsupervised learning are often episodic
 - E.g., speech recognition, medical diagnosis, credit score prediction, ...
- Reinforcement learning is often sequential
 - E.g., game playing

1. UNSUPERVISED LEARNING

- Given a set of objects: X
- Density estimation, $P(X)$
 - E.g., Bayesian networks
- Clustering
 - E.g., K-Means, Expectation Maximization (EM), Latent Dirichlet allocation (LDA), DBScan, ...
- Dimensionality reduction
 - Principal component analysis (PCA), independent component analysis (ICA), ...


2. SUPERVISED LEARNING

- Given objects with their labels, $\langle X, Y \rangle$
- Learn a function f that maps objects, X , to labels, Y
- We want f to perform well on unseen objects
- Several applications
 - Face recognition, speech recognition, medical diagnosis, fraud detection, credit scoring, home value prediction, temperature prediction, ...
- If Y is
 - Discrete, the task is called classification
 - Continuous, the task is called regression


FUNCTION FITTING?

- Isn't classification/regression simply “function fitting?”
 - Yes and No
- The purpose is to generalize and perform well on unseen data
- We don't want to **underfit** or **overfit** to the training data


CURVE FITTING




CURVE FITTING




CURVE FITTING




CURVE FITTING



CURVE FITTING



CURVE FITTING





So, which function is the "right" one?

CLASSIFICATION

- Decision trees
- Nearest neighbors
- Naïve Bayes
- Logistic regression
 - Note: it's called regression, but it is a classification model
- Support vector machines
- Neural networks

DECISION TREES



Learning: how do you learn a small tree that generalizes to unseen data? 17

NAÏVE BAYES

- Given X_1, X_2, \dots, X_n , and class Y
- Assume $X_i \perp X_j \mid Y$

$$P(Y|X_1, X_2, \dots, X_n) = \frac{P(X_1, X_2, \dots, X_n|Y)P(Y)}{P(X_1, X_2, \dots, X_n)} = \frac{P(Y) \prod_{i=1}^n P(X_i|Y)}{P(X_1, X_2, \dots, X_n)}$$

We need to estimate $P(Y)$ and $P(X_i | Y)$

LOGISTIC REGRESSION


- Learns $P(Y|\mathbf{X})$ directly, without going through $P(\mathbf{X}|Y)$ and $P(Y)$
- Assumes $P(Y|\mathbf{X})$ follows the logistic function

$$P(Y = \textit{false} \mid X_1, X_2, \dots, X_n) = \frac{1}{1 + e^{w_0 + \sum_{i=1}^n w_i X_i}}$$

$$P(Y = \textit{true} \mid X_1, X_2, \dots, X_n) = \frac{e^{w_0 + \sum_{i=1}^n w_i X_i}}{1 + e^{w_0 + \sum_{i=1}^n w_i X_i}}$$


- Learning: estimate the weights w_0, w_1, \dots, w_n

SUPPORT VECTOR MACHINES



$$\min \frac{1}{2} w^T w \text{ subject to } y^{(d)}(w^T x^{(d)} + b) \geq +1$$

NEURON



By Quasar Jarosz at English Wikipedia, CC BY-SA 3.0,
<https://commons.wikimedia.org/w/index.php?curid=7616130>

WHAT AN ARTIFICIAL NEURON DOES

- Takes a weighted sum of its inputs
 - $w_0 + \sum_{i=1}^k w_i x_i$
 - Assume that there is always a constant input 1, that is, $x_0 = 1$. Then,
 - $\sum_{i=0}^k w_i x_i$
- Passes this sum through its activation function
 - $f(\sum_{i=0}^k w_i x_i)$

MULTILAYER NEURAL NETWORKS

- An input layer
 - One or more hidden layers
 - An output layer
-
- Learning: estimate the weights

REGRESSION

- Linear regression
- Ridge
- Lasso
- Support vector machines
- Decision trees
- Nearest neighbors
- Gaussian processes

3. REINFORCEMENT LEARNING

- Agent interacts with the environment and receives rewards
 - E.g., play a game and receives a reward (which could be negative) after the game is over
- Goal: learn to maximize future rewards
- Might not receive a reward for each action
 - E.g., in the game of chess, the agent does not receive feedback for each move
 - Providing a reward at the end is easier to specify than providing feedback on each move
- Might need to learn the transition model (from one state to another) and/or the reward function
- Has to balance between exploration and exploitation

WE'LL COVER

- Unsupervised learning
 - Maximum likelihood estimation for density estimation in Bayesian networks
- Supervised learning
 - Naïve Bayes
- How related courses touch on ML
 - CS 584 – Machine Learning
 - All three types of ML; focuses on unsupervised and supervised learning
 - CS 577 – Deep learning
 - CS 422/522 – Data mining
 - Decision trees, random forests, support vector machines with a focus on data
 - CS 581 – Advanced AI
 - Reinforcement learning (and of course other advanced AI topics)

SCIKIT-LEARN

- An open-source Python package for primarily unsupervised and supervised learning
- <https://scikit-learn.org/stable/>
 - Preprocessing
 - Classification
 - Regression
 - Clustering
 - Dimensionality reduction